Roma Just Talk 元創盃組隊說明

給媽媽看的預覽｜2026 年 9 月 28 日｜尚未報名

媽，我想用正在開發的 Mac 語音輸入軟體 Roma Just Talk 參加元智大學第一屆「元創盃」創新創業競賽。現在最需要的是認識願意真正一起做提案、驗證產品與上台簡報的元智在學生。請幫我轉介有興趣的學生；先讓我和對方談工作內容，再決定是否組隊。

這是什麼產品

一般語音輸入要先按鍵、再開始說。Roma 的目標是讓人先說話，按下快捷鍵時也能保留前面短短幾秒的內容，減少漏掉句首與重講。已有公開的 Mac beta 和原始碼；目前仍需補足實際使用者驗證與比賽用的可重現展示。它是以 GPLv3 開源專案 VoiceInk 為基礎的 fork，我們會清楚說明既有成果與新增貢獻，不會說整個產品都由新隊伍從零做出。

組隊的硬條件

- 每隊 **3 至 6 位在學學生**，由 **1 位元智在學生擔任隊長**，而且全隊至少一半是元智學生。跨校學生可以加入，但要交在學證明。
- Felix 的學生身分與是否能列為隊員，需先確認。若不是在學學生，不能把 Felix 算進 3 至 6 人名額；仍可討論如何以既有專案負責人的身分協助學生隊伍，並向主辦單位確認權利與原創聲明。
- 不需公司、不需指導老師；但每位隊員應有真實分工並同意參賽。獎金分配與專案權利須先談清楚，不在邀請時預先承諾。

為什麼值得問學生

有可展示的產品，因此適合「實戰組」：初審看問題與市場 25%、解法與創新 25%、可行性 20%、已完成成果 20%、團隊分工 10%。學生能實際負責使用者訪談、Mac 展示與測試、提案、市場驗證或簡報。這些工作會形成可呈現的成果，不只是掛名。

時間： 2026 年 **10 月 16 日 17:00** 前完成線上報名與兩份 PDF；若晉級，**11 月 11 日**須參加提案診所，**12 月 7 日 17:00**前交決賽簡報，**12 月 11 日**到場做 5 分鐘簡報與 3 分鐘問答。10 月 7 日有自願參加的提案工作坊，公告與報名頁顯示的時間不同，出席前應向主辦單位確認。

可以轉給學生的短訊

> 元智同學你好，Felix 想以 Roma Just Talk 參加 2026 元創盃「實戰組」，正在找願意實際參與的學生隊友。Roma 是已有公開 beta 的 Mac 語音輸入專案，特色是「先說話、再按快捷鍵」也能保留句首。希望隊友協助使用者訪談、產品測試、提案或 12 月簡報；不要求每個人都會寫程式。比賽規定 3–6 位在學學生、元智學生隊長、至少一半元智學生，10/16 17:00 截止。專案來自開源 VoiceInk，我們會如實揭露來源與新貢獻。若你有興趣，請先聯絡 Felix 討論你想負責的部分與可投入時間：happyfumd@icloud.com。活動資訊：https://johnnyweb-art.github.io/yzu-startup-challenge/

給媽媽的下一步： 先找至少 3 位有興趣的元智在學生介紹給我；若 Felix 確認具在學學生資格，實際所需人數再依組隊比例調整。我來確認資格、分工、文件與主辦單位對開源 fork 的要求。請不要代替學生承諾參賽。

資料來源： 活動頁、正式簡章、提案格式。

YZU Startup Challenge 2026 — practical-track proposal draft

- **Team name:** [20 characters maximum; choose with team]
- **Proposal name:** Roma Just Talk
- **One sentence:** Speak before Shift.
- **Sector:** AI and ICT
- **Captain:** [YZU student, department/year]
- **Members and roles:** [3–6 total; use names and department/year only in proposal]
- **Advisor:** [optional]
- **Language:** English

> Export to selectable-text A4 PDF using the organizer's format: 12 pt Arial or Times New Roman, 1.5 spacing, numbered pages. Maximum eight body pages, separate cover; cite any appendix. Do not include IDs, phone numbers, or addresses. The name and one-liner must match the form exactly.

1. Summary

Roma Just Talk is a macOS dictation prototype for people who think of a sentence before they can reach a keyboard shortcut. Conventional push-to-talk starts recording after activation; a first word can be lost or the speaker has to pause. Roma keeps a short rolling microphone buffer on the Mac and joins it to audio captured after the shortcut. The intended interaction is to speak naturally, press Left Shift, then release to insert text into the active app. Users can choose local or configured cloud speech recognition. A public beta and source repository exist; the team has not yet established a verified customer count or repeat-use metric. Our next step is to test whether this interaction reduces lost opening words and interruption in real writing tasks, then learn whether users will adopt or pay for it. We seek users for a controlled pilot and advice on distribution, accessibility, privacy, and commercialization. The current app is derived from GPLv3 VoiceInk; our proposal concerns the differentiated interaction and continuing development, with third-party rights disclosed below.

2. Problem and market

Concrete moment. A student, writer, or knowledge worker is composing in a Mac app. They start speaking as an idea arrives but the dictation shortcut has not yet been pressed. Standard recording misses the beginning, so they repeat themselves, edit the first phrase, or stop to prepare the input tool. This friction is small per episode but interrupts a frequent workflow. Roma targets people who write across many desktop apps, especially those who already try voice input and are dissatisfied with activation delay.

First reachable segment. Begin with 10–20 Mac users who dictate multiple times a day, recruited through the existing public project and YZU/community channels. Run short task observations and a two-week opt-in pilot. Measure first-word capture, time to usable text, correction burden, and retention. This is a testable beachhead, **not** a claim about market size. Estimate TAM/SAM/SOM only after obtaining sourced platform and willingness-to-pay data; do not insert an invented revenue figure.

Current alternatives. Keyboard typing, built-in macOS dictation, and third-party dictation apps. Users may already be satisfied with these alternatives; the study must identify when Roma is meaningfully better and when it is not.

3. Solution and technology

The user speaks, presses Left Shift, continues speaking, and releases to insert text. Roma maintains a short pre-roll buffer so the opening phrase can enter the transcription pipeline. The app integrates with the macOS text-insertion flow and offers local model options or user-configured cloud providers. The distinctive product hypothesis is the **speak-before-hotkey** interaction, not a proprietary speech model.

The published codebase is a fork of VoiceInk. VoiceInk's macOS base and open-source speech libraries remain credited. Roma's team should itemize its own design and implementation contributions, especially pre-roll capture and the buffered-to-live handoff, using a commit-backed appendix. Distribution and reuse must respect GPLv3 and each dependency's license. **Organizer clarification on the declaration of originality remains pending.**

Privacy and reliability. Local transcription is an option. Cloud transcription sends data to the configured provider; it must be described accurately. The current public beta's signing, notarization, and data-handling limitations should be shown in the Trust disclosure, then checked against the release artifact before a public pitch. Do not imply a production-ready safety guarantee.

AI use disclosure. Speech recognition powers the product. Generative AI tools have also helped build and document the code, under human product and acceptance decisions. State the specific tools and scope used by the actual entrant team before submission. The contest permits AI tools and asks for truthful disclosure.

4. Existing work and validation

Built: a public source repository, v1.95.1 beta release, and a browser interaction demo. The browser demo illustrates the shortcut flow; it is not itself proof that the Mac app works end to end. Include dated screenshots of a verified app run and a working public demo link in the separate evidence PDF once checked.

Not yet validated for this submission: active users, retention, paid demand, first-word recovery rate in a user study, and measured advantage over alternatives. A downloaded release and automated build do not prove these outcomes. The pilot will use a fixed sentence/task set and record version, macOS version, hardware, microphone, model, exact spoken text, trigger timing, output, corrections, and consent. This makes the product claim falsifiable.

5. Business model

Hypothesis: a free evaluation path could lead to paid convenience or support for heavy users, or an organization could sponsor deployment. No price or plan is committed here. First interview target users about existing spend and test a concrete offer against alternatives. Because the application code is GPLv3, any paid model must respect the right to source and redistribute under that license. Do not claim exclusive ownership of the inherited code or promise a proprietary code license. Costs to estimate include model hosting when cloud options are used, release signing, support, and acquisition. The first distribution path is the public release and a clear onboarding page; effectiveness remains unmeasured.

6. Competition

Option	Useful strength	Question Roma must answer
Typing and editing	Familiar, precise, no mic setup	Is voice input faster for the target task?
Built-in macOS dictation	Present on the Mac, low acquisition friction	Is first-word recovery important enough to switch?
Established dictation apps	Mature recognition and workflows	Can Roma's interaction and control justify setup and trust costs?

The comparison should be tested on real tasks; no unsupported accuracy, RAM, price, or speed ranking is proposed. The differentiated shortcut can be copied, so durable value would need to come from execution quality, reliability, user trust, and workflow learning.

7. Team

Member	YZU or other school, department/year	Owned work	Evidence of ability
[Captain]	[YZU]	Product direction, submission, user research	[Specific work]
[Member 2]	[School]	macOS engineering and demo reliability	[Specific work]
[Member 3]	[School]	Market research, pilot recruitment, business model	[Specific work]

Assign actual work and confirm each person's consent. An advisor is optional. If a role is unfilled, name the gap rather than listing a nominal teammate. The product predates this competition; distinguish founder work and new teammate contributions honestly.

8. Next six months and resources

By	Deliverable	Completion test
2026-10-16	Eligible team and application	Center confirms received and then accepted; both PDFs conform
2026-11	First-word pilot	10+ consented Mac users; versioned task records and correction results
2026-12	Final pitch and live demo	Rehearsed 5-minute talk; 3-minute Q&A; offline video backup
2027-01	Onboarding and release quality	Tested published download on target macOS versions, documented blockers

By	Deliverable	Completion test
2027-03	Demand decision	Repeat-use and willingness-to-pay evidence; explicit continue/pivot choice

Requested support: introductions to target users and a mentor for Mac software distribution or open-source commercialization. Budget to be filled only after team agrees on line items and quotes: participant incentives, test hardware/access, Developer ID/notarization, and hosting if needed. Consider FITI/U-start with center guidance after eligibility and IP questions are resolved.

Source and rights note for appendix

VoiceInk: original upstream macOS application, GPLv3. Roma Just Talk: a public fork that retains VoiceInk credits and adds a speech-first pre-roll direction. Third-party libraries and model licenses should be listed from the release's dependency manifest. The team's own contribution list must be verified against Git history before the captain attests to originality. The organizer's rules say third-party technology and licenses must be stated; the captain should seek a written ruling on this fork before submission.

Roma Just Talk existing results

YZU Startup Challenge 2026 practical track evidence preview

Draft for family and prospective teammates. Not yet a submission. Team name, verified Mac capture, and participant results remain open. Export separately from the proposal as 元創盃_B_<team>_成果佐證.pdf after these are filled.

1. Public prototype

The Roma Just Talk source and v1.95.1 release are public. The release page published a Mac app archive on 2026-09-11. This shows that the project has progressed beyond an idea, but an archive is not proof of the exact live interaction under contest conditions.

Exhibit A to add: Two dated screenshots from the *published* Mac app: (1) shortcut/recording state, (2) text inserted into a real target app. Include the spoken phrase, exact output, macOS version, Mac model, microphone, speech model, and archive SHA/checksum. If a development candidate is used instead, identify it separately and do not label it the public release.

2. Demonstration

The browser demo source illustrates the speech-first interaction. It uses browser speech recognition and is distinct from the Mac app. A public URL should only be listed after opening it signed out, granting microphone access, and checking the interaction. Keep that URL reachable through at least 2026-12-11.

Exhibit B to add: A short recording of an actual Mac run with a visible pre-trigger phrase and final inserted text. Test the exact file or URL judges will receive. Keep a local playback copy for the December pitch.

3. Contribution and rights record

Roma is based on VoiceInk, which is GPLv3 licensed. Its inherited Mac base and dependencies must be credited. The entrant team should list its own pre-roll interaction, implementation, testing, design, and market-validation contributions with commit links and dates. It must not represent upstream code or pre-existing founder work as newly created by recruited students.

Exhibit C to add: 3–5 commit or design links and a table with author, date, contribution, and whether the author is on the eligible student team. Obtain the organizer's guidance on how to declare the fork and pre-existing work before the captain signs the originality statement.

4. External validation

No active-user count, paid-customer count, retention rate, or measured speed advantage is established in this preview. If pilot results arrive before submission, add the exact date range, sample, task protocol, raw counts, and anonymized findings. If there are no results, say so and present a concrete pilot plan rather than invented traction.

Exhibit D to add, if available: A compact task table showing first-word capture and correction time compared with the participant's current method. Link to a consented, anonymized result record.

Submission check

Single selectable-text PDF; all links accessible without login; screenshots legible; no private student data; under 100 MB (20 MB recommended). Confirm the team and rights statements match the proposal and form.